

EECS498-008

Formal Verification

of Systems Software

Material and slides created by
Jon Howell and Manos Kapritsos

About me

Manos Kapritsos (manosk@umich.edu)


Areas of research: Formal Verification, Distributed Systems

28 years old (well, maybe not in decimal)

Disclaimer

Any resemblance to purple-looking super villains, living or dead, is purely coincidental

About me



About you and me

- I **love teaching** and interacting with my students
- I want to get to know you all by name
 - **Send me your picture** by email (manosk@umich.edu)
 - Use subject "**EECS498-008 picture**"
- I'm here to help. Come to me with any question!
 - course-related: office hours (Tuesday, 11am-12pm, BBB 4824)
 - Life, The Universe, and Everything: any time

If you need more help...

Our TA, Armin Vakil, is here to help



A Jedi Master in Formal Verification

Office hours: 2-3pm Thursday, Learning Center, Table 2

The unseen hero

Jon Howell, VMWare Research



“Bugs, unit tests, gdb. The dark side are they.”

“Testing not make one great.”

Co-designer of the material in this course

Agenda for Today

- Why learn formal verification?
- Course syllabus and logistics
- What is formal verification?
 - (and other, similar approaches)
- Getting started with Dafny

Why learn formal verification?

EECS280/281


High-level programming

EECS482

OS-level programming

EECS491

Distributed programming



We have taught you how to write and test programs

Real-world systems are too complex to test




Figure 1: Typical Figure 2 from Byzantine fault paper: Our network protocol



[Mickens 2013]

Despite tremendous effort...



“not one of the properties claimed invariant in [PODC] is actually invariantly true of it.”



[Zave 2015]

Under the same assumptions made in the Chord papers, the [SIGCOMM] version of the protocol is not correct, and not one of the properties claimed invariant in [PODC] is actually invariantly true of it. The [PODC] version satisfies one invariant, but is still not correct. The results are presented by means of counterexamples to the invariants in Section 4. In preparation for the results, Section 2 gives a brief summary of the protocol and failure assumptions, and Section 3 introduces the invariants.

Buggy code update triggers Bing, Yahoo outage

Jan 5, 2015 9:19 AM
Filed under [Software](#)

Like 9 Tweet 14 +1 0 in Share 4



Siri and Cortana also affect

Yahoo and Microsoft search engine for 20 minutes after Microsoft pushed a software update that struggled to roll it back.

Tags: Microsoft's Bing search engine powered by Microsoft announced in 2009.

Google blames software bug for Friday night Gmail outage



Promoted Content

Innovation at work
Want to find out how to unleash innovation in your organisation? Whether you're keen to learn more about hotdesking, laser printing, connectivity, mobility, security or more, check out our...

10 Things Your Next Firewall Must Do
Start Thinking: Next-generation firewall.

Search giant apologises for outage that left millions of users unable to access services.

WRITTEN BY Caroline Donnelly

Is it possible to write code that is
completely bug-free?

Formal verification is as close as we can get

March 25, 2015

(or how I became a believer)

- One day before the deadline of SOSp'15, we had not yet run our code
- ...and yet our code **ran correctly the first time** we ran it!
 - (and every time afterwards)

How will you benefit from this class?

- It will make you a better programmer
 - Whether you end up writing verified code or not
- You will learn to specify your code
 - To express your intent clearly and unambiguously
 - Just like design docs, but better
- You will learn important concepts
 - E.g. inductive invariants, refinement
- Get ahead of the curve
 - Learn an emerging skill

Objectives of this class


- Understand the **fundamentals** of formal verification
- Learn **how they apply** in a (distributed) systems context
- Get **hands-on experience** with proving systems correct
- Become familiar with a **practical verification language**

Prerequisites

- Experience with **distributed** or **asynchronous** systems
 - i.e. EECS482
- I will explain the mechanics of any distributed systems in the class for those of you that haven't taken EECS491
- *No verification experience required*

If **anything** is unclear, do not hesitate to ask

About this class

- Disclaimer: this class is not formally verified! 
- This is the first time I am teaching this class
 - ...or rather, anyone is teaching this class
 - ...or rather, anyone is teaching a class on this material
- There is no textbook (anywhere!)
 - Jon and I will write one based on the experience from this class

Class material

- Class webpage
 - <https://verification.eecs.umich.edu>
- Syllabus, lecture slides, problem sets and projects will be posted on the class webpage
- **Subscribe yourself to Piazza**
 - Announcements and class discussion

Enrollment

- Current cap at 40 with a few people on the waitlist
- I'm working to increase the cap (and perhaps add a second lab)

Lectures

- Lectures will be held in person in DOW 1017
- Recordings will be posted at:
<http://leccap.engin.umich.edu/leccap/site/p3kt6ubmxl1hk61vfm0>
(link is also on course web page and Piazza)
- I will be posting slides of each lecture (shortly) before the lecture, in case you want to keep notes directly on the slides

Lectures schedule

- Material is divided into eight chapters
 - Chapters 1-5 (before midterm) cover the basic concepts of verification
 - Basic verification and Dafny mechanics
 - Specification
 - Centralized state machines
 - Proving properties and inductive invariants
 - Distributed state machines
 - Chapters 6-8 (after midterm) cover refinement proofs and advanced topics
 - Refinement
 - Asynchronous clients
 - Application correspondence

Problem sets and projects

- There are four (programming) problem sets and two projects
 - Problem sets will be done individually
 - Projects will be done in groups of 1-2 students
- All deliverables will be submitted via the autograder.io website
 - They use a combination of auto-grading and hand-grading

(this is where I warn you about how hard the projects are, but...)

Policies

- Submission
 - Three submissions per day to the autograder
 - Due at midnight on deadline
 - Three late days throughout the semester
- Collaboration
 - Okay to clarify problem or discuss Dafny syntax
 - Not okay to discuss solutions



Exams

- Midterm: October 12, 6-8pm
- Final: December 12, 10:30am-12:30pm
- No makeup exams
 - Except in dire circumstances
 - Make sure you schedule your interviews appropriately

Grading breakdown

- Problem sets: 30%
 - PS1: 8%
 - PS2: 8%
 - PS3: 8%
 - PS4: 6%
- Projects: 30%
 - Project 1: 15%
 - Project 2: 15%
- Midterm exam: 20%
- Final exam: 20%

What is formal verification?

- Step 1: **Specify** the correctness of the system formally
- Step 2: **Prove** that the implementation conforms to the spec

If the spec expresses your correctness property,
then your system is **correct**, subject to any
assumptions you have made during your proof

Other approaches

Testing: run the system with a large and/or representative set of inputs to determine if it behaves correctly

- **Quality depends** on acumen of test designer
- **Infeasible** to achieve complete coverage for complex systems

Model checking: Model the system and ensure all possible states are safe

- **Correctness depends** on how accurate the model is
- **Does not scale** well to complex systems, especially those with infinite state spaces, like distributed systems

Statically checking for correctness

What we want is a “static correctness check”, akin to a static type check

You write your code normally, but if you introduce bugs the checker will tell you

When the checker complains, you have to spend some time to convince it that your code is right---if indeed it is

Using a Theorem Prover

Express the execution of the system and its correctness as a **mathematical formula** (done automatically by the language)

Give the formula to a theorem prover, effectively asking:
"If the system behaves this way, is it possible for its correctness to be violated?"

A **negative** answer means the system is **proven to be correct**

A **positive** answer means there is still work to do, either:

- the system is indeed incorrect
- the proof is incomplete



Using Dafny

- We will be using Dafny as our verification language
- Dafny is an **imperative** language designed with formal verification in mind
 - ...and plenty of functional language features
- Dafny uses an SMT solver (Z3) to automate verification to a large degree
 - ...but it needs our help sometimes
- Most of the high-level skills are transferrable...
 - ...but some are specific to Dafny and/or automation

Dafny in Docker



- We provide you with a Docker container that has Dafny pre-installed
 - Makes it easy to get started
 - Ensures everyone is using the same Dafny version as the autograder
- Download and run it like this:
 - `docker pull ekaprits/eecs498-008`
 - `docker container run --mount src=$PWD,target=/home/autograder/working_dir,type=bind,readonly -t -i ekaprits/eecs498-008`
- CAEN machines don't support Docker, yet
 - If you don't have access to a machine that can run Docker, contact me ASAP

Things to do

- Browse the course web page
- Subscribe to Piazza
- Install Docker, pull image, run image