

EECS498-008

Formal Verification

of Systems Software

Material and slides created by
Jon Howell and Manos Kapritsos



PREVIOUSLY ON
FORMAL VERIFICATION



X Slides

M | COMPUTER SCIENCE & ENGINEERING

Defining the network

Distributed system

10/5/22

```
datatype Option<T> = Some(value:T) | None
datatype MessageOps = MessageOps(
  recv:Option<Message>,
  send:Option<Message>)
```


Network module

```
module Network {
  datatype Variables =
    Variables(sentMsgs: set<Message>)

  predicate Next(v, v', msgOps:MessageOps) {
    // can only receive messages that have been sent
    && (msgOps.recv.Some? ==> msgOps.recv.value in v.sentMsgs)
    // Record the sent message, if there was one
    && v'.sentMsgs ==
      v.sentMsgs + if msgOps.send.None? then {}
                  else {msgOps.send.value}
  }
}
```

EECS498-008 11

Defining the network



Network module

```

module Network {
  datatype Variables =
    Variables(sentMsgs: set<Message>)

  predicate Next(v, v', msgOps:MessageOps) {
    // can only receive messages that have been sent
    && (msgOps.recv.Some? ==> msgOps.recv.value in
      v.sentMsgs)

    // Record the sent message, if there was one
    && v'.sentMsgs ==
      v.sentMsgs + if msgOps.send.None? then {}
                  else {msgOps.send.value}
  }
}
  
```


```

datatype Option<T> = Some(value:T) | None
datatype MessageOps = MessageOps(
  recv:Option<Message>,
  send:Option<Message>)
  
```

A distributed system is composed of multiple hosts and a network

Distributed system: attempt #2

```
module DistributedSystem {
  datatype Variables =
    Variables(hosts:seq<Host.Variables>,
              network: Network.Variables)
```



```

predicate HostAction(v, v', hostid, msgOps) {
  && Host.Next(v.hosts[hostid], v'.hosts[hostid], msgOps)
  && forall otherHost:nat | otherHost != hostid ::
    v'.hosts[otherHost] == v.hosts[otherHost]
}

predicate Next(v, v', hostid, msgOps: MessageOps) {
  && HostAction(v, v', hostid, msgOps)
  && Network.Next(v, v', msgOps)
}
  
```

Binding variable

Administrivia

- Midterm exam **this Wednesday, 10/12**
 - 6-8pm, EECS1303
 - No lecture that day
- Closed books
 - Allowed one double-sided “cheat-sheet”, 10pt minimum
- Covers everything up to Chapter 4 (i.e. excluding distributed systems)
- Problem set 3 (Chapter 5) will be released later today
- Start looking for partners for Project 1 (released after PS3)

Atomic Commit (Problem Set 3)



- Do you take each other?
- I do.
- I do.
- I now pronounce you
atomically committed.

Atomic Commit: the objective

Preserve data consistency for distributed transactions

Example: book a hotel and flight on Expedia

Atomic Commit: the setup

- One coordinator
- A set of participants
 - Allowed to be empty in our model
- Every participant has an “input” value, called **vote/preference**
 $vote_i \in \{Yes, No\}$
- Every participant/coordinator has an “output” value, called **decision**
 $decision_i \in \{Commit, Abort\}$
- We are ignoring the possibility of failures

Atomic Commit: the spec (simplified to ignore failures)

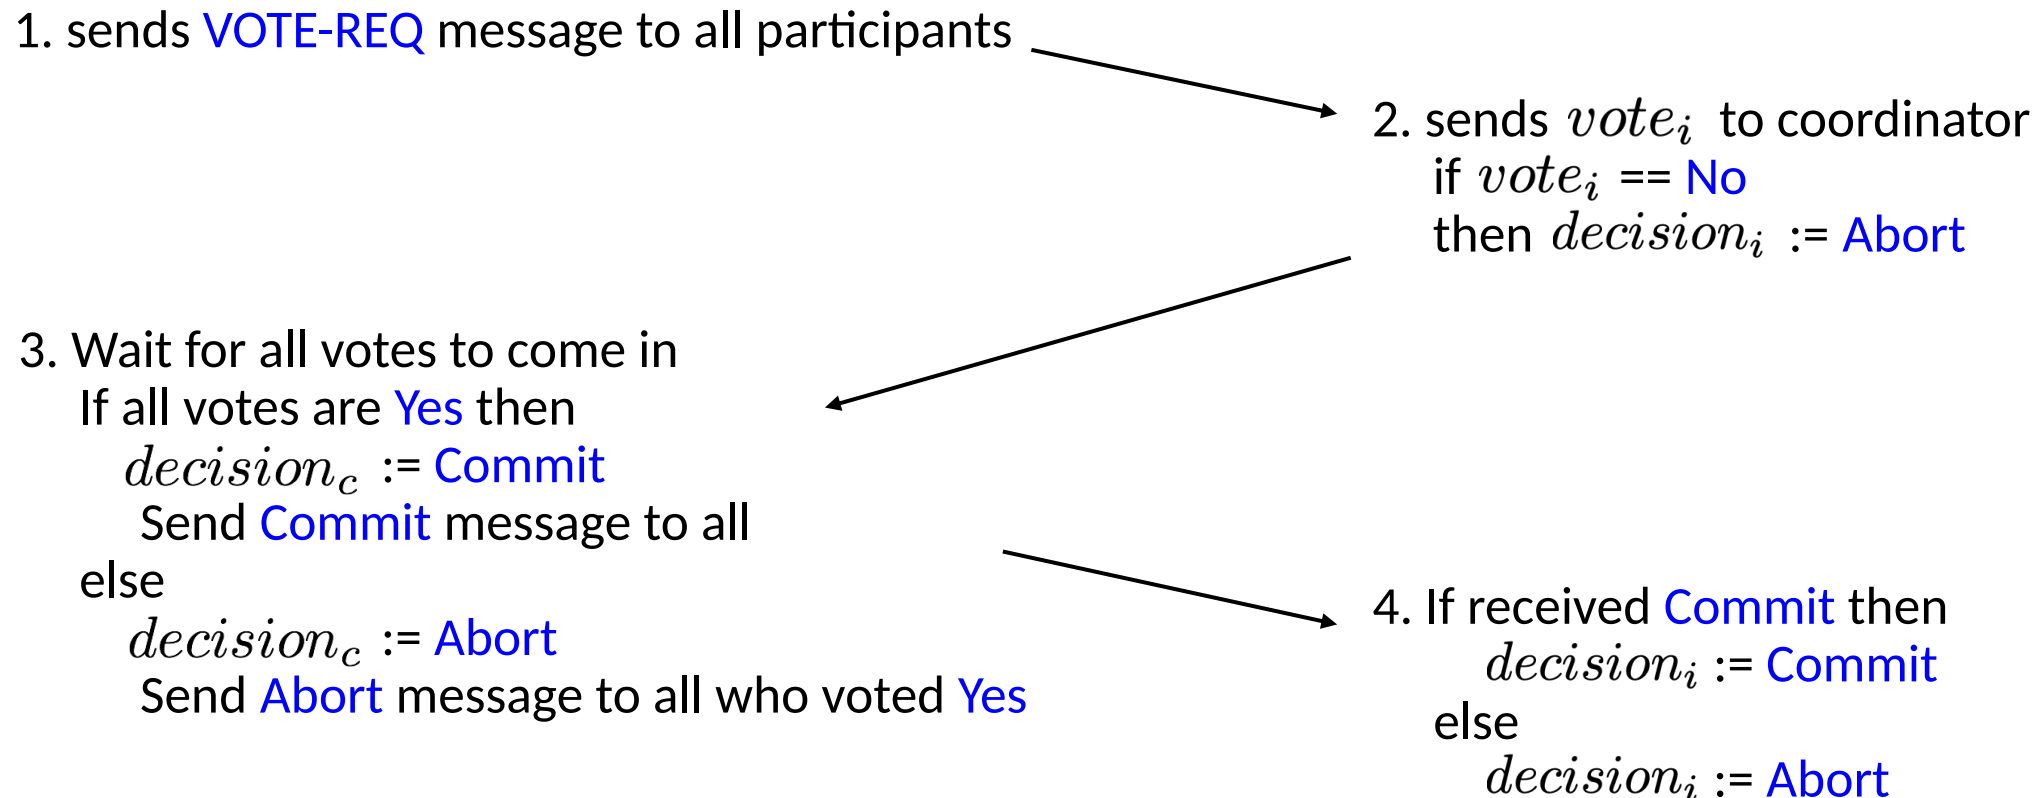
- AC-1: All processes that reach a decision reach the same one
- AC-3: The **Commit** decision can only be reached if all processes vote **Yes**
- AC-4: If ~~there are no failures and~~ all processes vote **Yes**, then the decision must be **Commit**

AC-2 and AC-5 ignored

Two Phase Commit (2PC)

Coordinator c

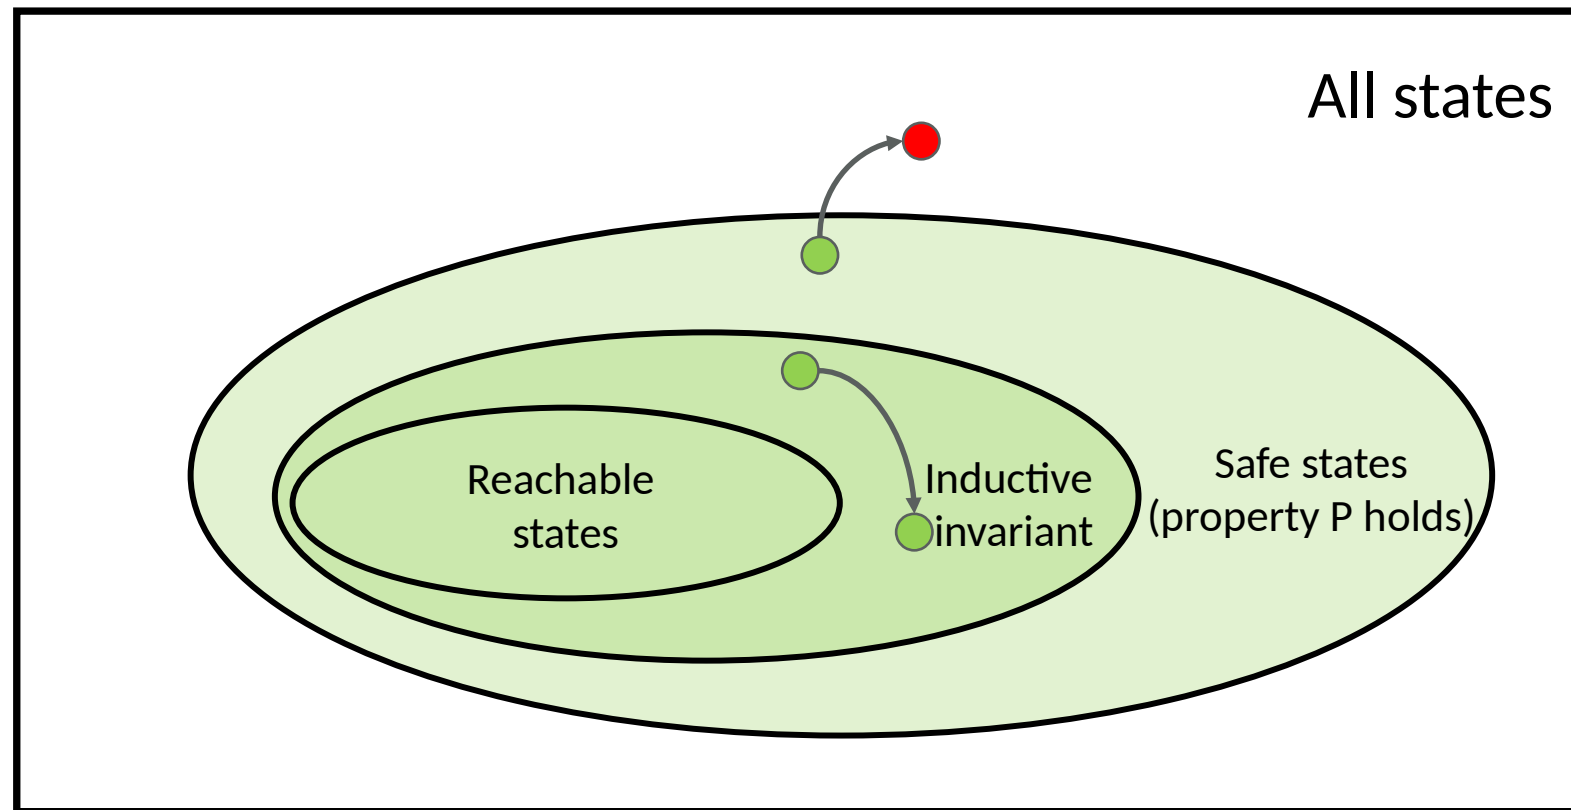
Participant p_i



Recap of Chapters 1-4

- Chapter 1: Dafny mechanics
 - Primitive types, quantifiers, assertions, recursion, loop invariants, datatypes
- Chapter 2: Specification
 - Formally define how a system should behave
- Chapter 3: State machines
 - Express the behavior of a system using Init() and Next() predicates, JNF
- Chapter 4: Inductive invariants
 - A strengthening of the safety property to become inductive

Invariants vs Inductive invariants



A distributed system is composed of multiple hosts, a **network** and **clocks**

Distributed system: attempt #3

```
module DistributedSystem {
```

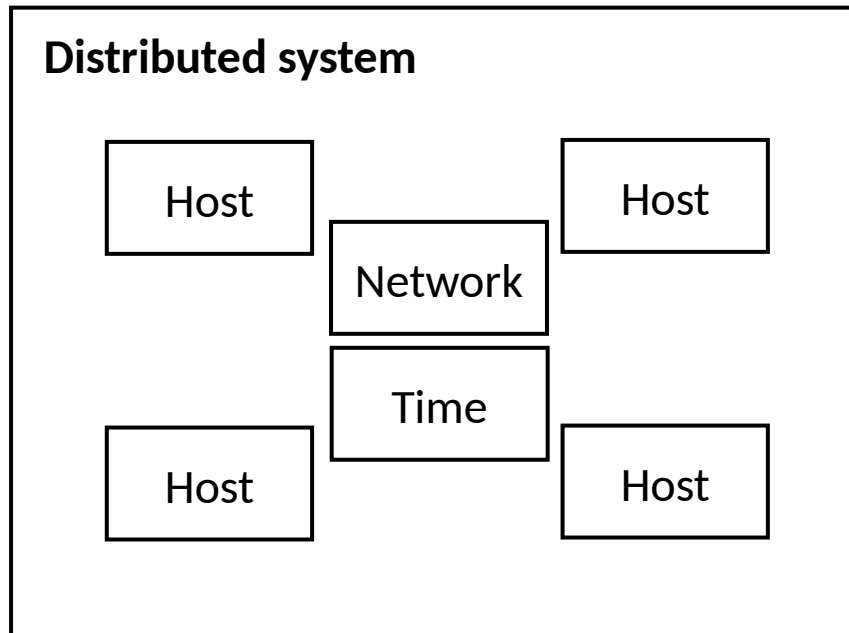
```
  datatype Variables =
```

```
    Variables(hosts:seq<Host.Variables>,
              network: Network.Variables,
              time: Time.Variables)
```

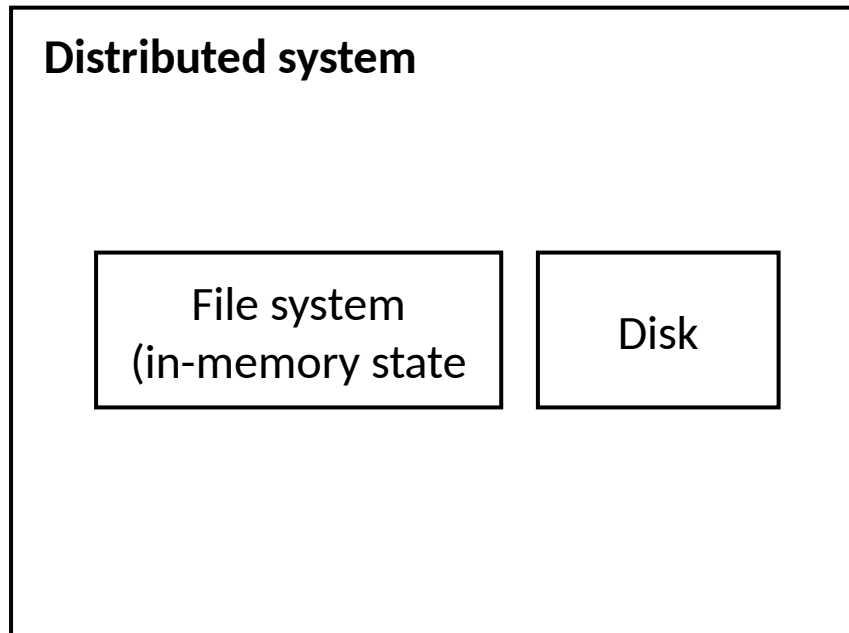
```
  predicate Next(v, v', hostid, msgOps: MessageOps, clk:Time) {
    || (&& HostAction(v, v', hostid, msgOps)
        && Network.Next(v, v', msgOps)
        && Time.Read(v.time, clk))
    || (&& Time.Advance(v.time, v'.time)
        && v'.hosts == v.hosts
        && v'.network == v.network)
```

```
  }
```

```
}
```



A “distributed” system



```



module DistributedSystem {
  datatype Variables =
    Variables(fs: FileSystem.Variables,
              disk: Disk.Variables)

  predicate Next(v, v') {
    || (exists io ::
        && FileSystem.Next(v.fs, v'.fs, io)
        && Disk.Next(v.disk, v'.disk, io)
    || ( // Crash!
        && FileSystem.Init(v'.fs)
        && v'.disk == v.disk
    )
  }
}

```

Binding variable

Trusted vs proven



SPECIFICATION : the systems specification sandwich



image: pixabay

